

Assignment 2: Total Recall

Surekha Gurung (sg1089)

Krishna Patel (kap335)

Baljit Kaur (bk404)

Malloc & Free

void* malloc (size_t size)

The user calls malloc to request a memory allocation of a certain size. Our header file defines malloc to perform the function myallocate(....., THREADREQ) which lets myallocate know that it is a memory request from user code. We also have a LIBRARYREQ to indicate a memory request from our library. Myallocate(...) performs different actions depending on the callers. We use first fit to find a block with enough space for the allocation. If there is no fit, myallocate(...) prints an error statement. For a THREADREQ call, myallocate(...) gets the threadID, initializes the pages, and searches for a free block with enough space for allocation (first fit) through the helper methods manage(), doReq, doPages, and doStuff(). It then updates the metadata of blocks by using bit shifts. Myallocate(...) also initializes the blocks, swapfile, signals, and pages before performing the different actions using the helper methods initblock(), setDisk(), and setPages().

free(void* address)

The user calls mydeallocate function to free a particular block of memory. The address of the block is passed and this function checks if the address is actually allocated or not. It also checks if the address is already freed or not. This function checks for both LIBRARYREQ and THREADREQ allocated blocks of memory. After freeing a particular block of memory, with a helper function called freeBlock() it merges with other free block of memory if there are any free blocks available.

Metadata

We stored our metadata in an unsigned int to make our program more memory efficient. We saved space by using only 4 bytes versus using around 10 bytes if we had used structs. The first 2 bytes represent the identification number of the page. The third byte is used to store the ID of the thread, and the last byte accounts for whether the page is occupied or free. We also used an unsigned int to store the metadata for blocks. The first three bytes are used to store the size of the block. Four bits are used to account for the vacancy of the block. Another four bits are used to store whether there is a block after the current block. If there is, to access it, we increment the address of the current location by its size. If there is no next block, we return null. We increment the addresses by bit-shifting in the functions themselves.

Pthreads

int my_pthread_create (...)

The my_pthread_create takes my_pthread_t pointer, an attr pointer which is set to null, function pointer of type void* and a void* argument. First, it checks if it is initialized or not. If not initialized, then it initializes a list of nodes, the run and promo queues, and creates a context. The create function returns 0 if successfully done.

int my_pthread_yield ()

In this function, no argument is passed. It first checks if the list is initialized or not. Then, the timer is set for all the threads created by the users. The action of all the threads is set to yielding. After, the scheduler is called to put the threads together in the priority queue.

void my_pthread_exit (void *value_ptr)

This function takes a void pointer argument and checks for initialization. If the initialization is done, the timer is stopped and the status is changed to exiting thread. Then, the scheduler is called and the threads are enqueued or dequeued in the priority queue.

int my_pthread_join (my_pthread_t thread, void **value_ptr)

The join function takes the thread and void** pointer as argument. First initialization is checked for the list. Then the thread is checked if it is dead or not, if not then threads are moved to the temp nodes and resumes the timer. If not dead then the thread's status is set to joining and then temp list is incremented.

Mutexes

int my_pthread_mutex_init (...)

This function takes argument as pointer to a mutex that is created by the user. In addition, it also takes argument as pointer to mutexattr that is set to null. First it checks for the mutex. If not null, the mutex lock is initialized to the threads and the flag is kept for the track of the locks.

int my_pthread_mutex_lock (my_pthread_mutex_t *mutex)

This function takes an argument as pointer to the mutex that is created by the users. Then after getting the mutex lock it checks the test and set condition and calls my_pthread_yield () for single thread and the yielding of threads are done.

int my_pthread_mutex_unlock (my_pthread_mutex_t *mutex)

This function takes an argument as pointer to the mutex that is created by the users. After the work is done the lock is unlocked and the flag is dropped to 0.

int my_pthread_mutex_destroy (my_pthread_mutex_t *mutex)

This function takes an argument as pointer to the mutex that is created by the users. It checks if the flag is dropped to 0 or not. If it is then the lock is destroyed and returns 0.

Scheduler

The scheduler struct has a pointer that points to the current node that is running. The node struct contains a pointer to a specific thread, its context, the joining threads, priority, and other information about that particular thread. We also included a list of all threads as well as a list of threads waiting for others to finish. This makes it easier to implement join and exit. The scheduler struct also has a multilevel priority queue of 5 levels. Each thread runs a time slice based on its level. Each level gets to run its threads for 50ms(priority level+1) ms in our implementation. As advised in the project description, threads are moved down to the next level queue if it runs for the full time.

The scheduler function has conditionals based on which on the thread's actions (joining, Exiting, yielding). After going through each level and doing the joining, exiting and yielding it gives back the threads. It also keeps tracks of the time taken and the cycles used by each thread. To avoid starvation, we used a check() function that is called every hundredth call to the scheduler. The check() function increases the priority of threads that are in the lower levels based on when it was created.

Testing

We tested our assignment using the Benchmarks provided by the TA. Our results for parallelCal.c matched the TA's results. For instance, running ./parallelCal 4 gives
sum is: 1726371200
verified sum is: 1726371200

In addition, ./vectorMultiply 4 gives
Res is: 332833500
Verified res is: 332833500

Hence, it is safe to say that our implementation gives the correct results, and it seems to be working well.

We also calculated that the runtime of our implementation if we ran the following:

[-bash-4.2\$./parallelCal 13456] is 0.76ms.
[-bash-4.2\$./vectorMultiply 13456] is 0.60ms.